
Music Creation through Non-negative Matrix Factorization

Jaskirat Kaur

Department of Computer Science
The University of Chicago
Chicago, IL
jaskirat@uchicago.edu

Sarang Joshi

Department of Computer Science
The University of Chicago
Chicago, IL
sarangj@uchicago.edu

Trenton Wesley

Department of Computer Science
The University of Chicago
Chicago, IL
trentonjw@uchicago.edu

Abstract

1 In this paper, we explore non-negative matrix factorization (NMF) in the context of
2 music samples. We show that NMF is a useful method for capturing the features
3 of audio files. We also show that different loss functions in NMF algorithms can
4 lead to a significant difference in the features captured. Finally, we successfully
5 combine the features of different music samples to generate new music.

1 Introduction

7 Non-negative matrix factorization (NMF) refers to a group of linear algebra techniques that allows
8 one to decompose a matrix into a product of two matrices. With NMF, a matrix, V is typically
9 factorized into two matrices, W and H . Commonly, W is referred to as the feature matrix and H
10 referred to as the coefficients matrix. All three of these matrices have a non-negativity constraint
11 that requires all elements to be greater than or equal to 0. This non-negativity constraint can be very
12 useful for interpretability. Overall, NMF is a handy technique that has applications in bioinformatics
13 [12, 7], computer vision [3], and astronomy [1, 13], among many others.

14 Previous work shows how NMF can be used to decompose music samples and derive representations
15 for the musical components of the sample and their temporal arrangement [5, 8, 9]. NMF has been
16 used in this way to remove certain sources of noise in music [6], to categorize music samples [9],
17 and to create new music through the combining of features from different music samples [2, 11].
18 NMF has proved itself as a computationally inexpensive method for analyzing audio. However,
19 there are a number of NMF algorithms that have their various uses in different areas. These algorithms
20 use different loss functions and can reach varying results. It remains an open question as to what
21 algorithms work best regarding audio files.

22 In this paper, we set out to explore various NMF algorithms typically utilized in a range of applications.
23 We will characterize these algorithms and their use in music decomposition, seeing how their results
24 differ.

25 2 Related Work

26 2.1 Music Decomposition

The basic method in which people decompose these music samples takes two steps. In the first step, we take an music input signal and run a Short-Time Fourier Transform(STFT) on it. This input signal is some floating point time series data representing a music file. Running STFT, we get a complex-valued spectrogram matrix D from which we can decompose using NMF for our second step. We can then separate this matrix into its magnitude S and phase P components. In doing so, we get back two matrices W and H , where

$$S = WH$$

27 Here, if $S \in \mathbb{R}^{n \times p}$, then $W \in \mathbb{R}^{n \times k}$ and $H \in \mathbb{R}^{k \times p}$ for some $k > 0$. U would be our feature
28 matrix, with each column u_i representing the individual musical components in our sample. W would
29 represent the temporal arrangement of the musical elements in U . Multiplying these two matrices
30 back together would give back a reconstructed version of our original matrix. From here, we can add
31 our P phase components and run an inverse STFT to get back an audio sample to listen to.

32 2.2 Cost Function

33 We consider two cost functions to assess the quality of our approximation of the NMF Factorization
34 $V \approx WH$. First we use Euclidean distance[10] which is widely used in measuring the dissimilarity
35 between two matrices given by

$$\|A - B\|^2 = \sum_{ij} (A_{ij} - B_{ij})^2$$

36 The minimum value of this measure is zero and it only reaches zero when A exactly equals B .
37 Another insightful matrix is defined as the Kullback-Liebler divergence:

$$D(A \| B) = \sum A_{ij} \log \frac{A_{ij}}{B_{ij}} - A_{ij} + B_{ij}$$

38 Similar to the Euclidean distance, the metric has a lower limit of zero and reaches zero only when A
39 is precisely equal to B . However, it lacks symmetry in A and B leading us to term it the "divergence"
40 of A from B . This matrix takes the form of the Kullback-Liebler divergence or relative entropy when
41 $\sum A_{ij} = \sum B_{ij} = 1$.

42 Finally, as a reference, we perform the NMF of the sample matrix using the `scikit-learn` library's
43 implementation [14] as a baseline.

44 2.3 Optimization for Music Decomposition

45 In our investigation, we address the music decomposition problem using non-negative matrix factor-
46 ization (NMF), which is given by the following optimization problems:[10].

$$\begin{aligned} \text{Problem 1: } & \min_{W, H \geq 0} \|V - WH\|_F^2 \\ \text{Problem 2: } & \min_{W, H \geq 0} D(V \| WH) \end{aligned}$$

47 Note that while both the objective functions are convex in either W or H , they are not jointly convex
48 in both variables. Consequently, finding a global optimum is not feasible. We can use gradient descent
49 to approach the problem, however, it has a slow convergence. We are instead using multiplicative
50 update rules to strike a balance between convergence rate and ease of implementation for NMF.
51
52

53 2.4 Multiplicative update rules

54 **Theorem 1:** *Under the iterative update rules, the euclidean distance*

$$H_{ia} \leftarrow H_{ia} \cdot \frac{(W^\top V)_{ia}}{(W^\top WH)_{ia}}, \quad W_{ia} \leftarrow W_{ia} \cdot \frac{(VH^\top)_{ia}}{(WHH^\top)_{ia}}$$

55 This distance remains invariant if and only if matrices W and H are at a stationary point with respect
 56 to it.

57 **Theorem 2:** *Under the iterative update rules, the divergence is non-increasing.*

$$H_{ia} \leftarrow H_{ia} \cdot \frac{\sum_k W_{ka} V_{ik} / (WH)_{ik}}{\sum_k W_{ka}}, \quad W_{ia} \leftarrow W_{ia} \cdot \frac{\sum_k H_{ka} V_{ik} / (WH)_{ik}}{\sum_k H_{ka}}$$

58 This divergence measure becomes stationary if and only if the matrices W and H have reached a
 59 stationary point with regard to this measure.

60 3 Methods

61 In order to compare the performance and results of the various objective functions in NMF, we
 62 chose two samples from the Ballroom dataset [4]: **I Like It 2** and **Latin Jam 2**, extracted 10-second
 63 samples, and performed the NMF-based musical decomposition with each of the three approaches
 64 discussed above: `scikit-learn`'s state-of-the-art approach, Euclidean distance, and divergence.
 65 Below is our implementation in Python of optimizing for the latter two objective functions:

```
def euclidean_nmf(V, n_components, n_iter = 1000):
    n, p = V.shape
    W = np.random.random((n, n_components))
    H = np.random.random((n_components, p))

    for i in range(n_iter):
        H_new = np.multiply(H, np.divide(W.T @ V, W.T @ W @ H))
        W_new = np.multiply(W, np.divide(V @ H.T, W @ H @ H.T))

        H = H_new
        W = W_new

    return W, H
```

```
def divergence_nmf(V, n_components, n_iter = 1000):
    n, p = X.shape
    W = np.random.random((n, n_components))
    H = np.random.random((n_components, p))

    for i in range(n_iter):
        W_H = W @ H
        H_new = np.zeros(H.shape)
        for a in range(H.shape[0]):
            for b in range(H.shape[1]):
                H_new[a,b] = H[a,b] * \
                    np.sum(np.divide(np.multiply(W[:,a], V[:,b]), W_H[:,b])) / \
                    np.sum(W[:,a])

        W_new = np.zeros(W.shape)
        for a in range(W.shape[0]):
            for b in range(W.shape[1]):
                W_new[a,b] = W[a,b] * \
                    np.sum(np.divide(np.multiply(H[b], V[a]), W_H[a])) / \
                    np.sum(H[b])

        H = H_new
        W = W_new

    return W, H
```

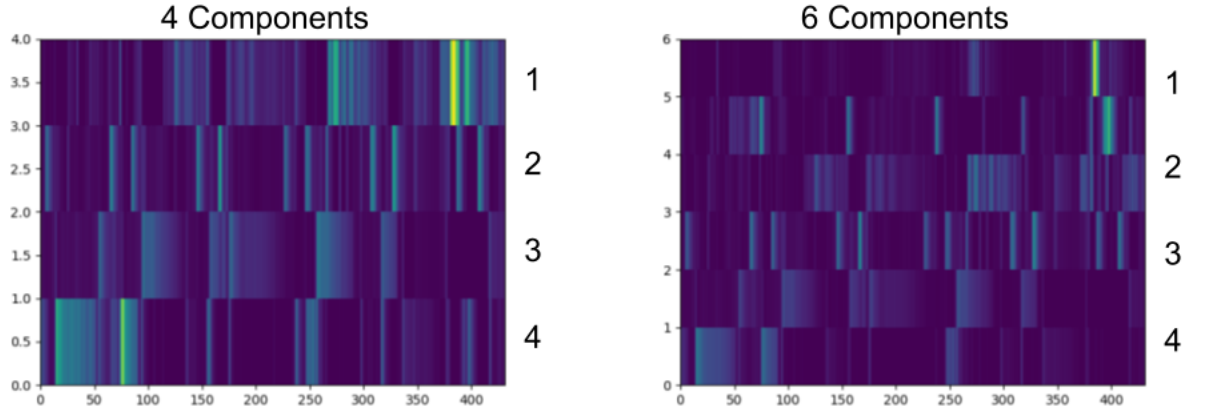


Figure 1: Shows H Matrices with 4 and 6 components respectively. Matrices were given by the use of the State-of-the-Art NMF decomposition method on **I Like It 2** audio file

First, we compared the output of running NMF with $n = 4$ and $n = 6$ components for the **I Like It 2** sample to contrast the component distribution of the two parameter values.

We chose $n = 4$ as the number of components in the two samples, based on the results of the previous comparison as well as a cursory auditory analysis of the range of instruments/frequencies in each.

4 Results & Discussion

4.1 Number of Components

We first observe the different H matrixes using the State-of-the-Art approach for **I Like It 2** with $n = 4$ and $n = 6$ components. We can see that the high value points of the components get translated/combined from $n = 6$ to $n = 4$. If we look at the concentrated high-value section in the top-right corner of our spectrograms, we can see that the $n = 4$ matrix's 4th row values have been spread across the 3rd and 4th rows of the $n = 6$ matrix at the same time sections.

As a result of this experiment, we decided to choose $n = 4$ as the number of components for the multi-sample comparison (in the next section). The components are more clearly defined with this value across the three approaches, and upon reconstruction of individual components, we see more clear differentiation with this value.

4.2 Feature Interpretation

The results of running the three NMF implementations are shown in the spectrograms for figures 2-5. For the W matrices, each column represents a component. For the H matrices, each row represents a component over the course of the 10-second sample. We can see from these figures that the different NMF approaches produced a different distribution of information across the 4 components.

In the State-of-the-Art approach for **Latin Jam 2**, we can see that the values of the 4th column has much higher values than the other three. This is somewhat similar to the Euclidean objective approach. However, in the divergence objective approach, the concentration of these high values is distributed much more evenly across the four columns.

Moving to the H matrices, which represent the temporal envelopes of the sample, we notice some strong similarities across the three approaches. We start with the graphs for **Latin Jam 2**. Take for




Figure 2: Shows W matrices for **Latin Jam 2**. Audio file is decomposed using the three compared methods: State-of-the-Art, Euclidean, and Divergence.




Figure 3: Shows H matrices for **Latin Jam 2**. Audio file is decomposed using the three compared methods: State-of-the-Art, Euclidean, and Divergence.

92 example the 3rd row in the State-of-the-Art approach. We see that there is a high value at a steady
 93 cadence, approximately every 30 columns. This appears to correspond to the steady beat that is in
 94 the music sample. Now if we look at the 1st row in the Euclidean graph and the 2nd column in the
 95 divergence graph, we see that the cadence is almost identical. The same similarity can be seen for **I**
 96 **Like It 2**: the 4th row in the State-of-the-Art approach loosely matches the 3rd row in the Euclidean
 97 approach and the divergence approach.

98 We can observe therefore that even though the exact nature of the individual components differs
 99 from approach to approach, all three are able to identify some fundamental musical aspects of the
 100 sample. The other components are also similar, with minor differences that are spread across other
 101 components.




Figure 4: Shows W matrices for **I Like It 2**. Audio file is decomposed using the three compared methods: State-of-the-Art, Euclidean, and Divergence.




Figure 5: Shows H matrices for **I Like It 2**. Audio file is decomposed using the three compared methods: State-of-the-Art, Euclidean, and Divergence.

5 Conclusion

In conclusion, we have investigated the application of non-negative matrix factorization (NMF) for music decomposition. We explored different NMF algorithms, each with different loss functions, that exhibited distinct results when applied to audio files. By applying Euclidean distance and the Kullack-Leibler divergence as cost functions, we were able to characterize the performance of these algorithms in reconstructing and interpreting musical components from a given sample.

Through careful analysis of the W and H matrices for our chosen music samples, it became apparent that while some algorithms, particularly the state-of-the-art approach from scikit-learn, lean towards isolating certain musical instruments, others, particularly those driven by the divergence objective, provide a more evenly distributed representation across components.

We also experimented with the versatility of NMF in music-mosaic, that is creating new musical compositions by amalgamating features from different sources. Hence, NMF is not only an analysis tool but can also be used as a creative platform for musicians and engineers.

6 Additional Ideas

Our exploration raises question of algorithm selection based on the genre, complexity and desired outcome of music analysis. Our future work can include exploring custom-tailored NMF solutions that cater to specific artistic requirements thereby opening possibilities for the domain of automated music production.

References

- [1] Olivier Berné, A Helens, Paolo Pilleri, and Christine Joblin. Non-negative matrix factorization pansharpening of hyperspectral data: An application to mid-infrared astronomy. In *2010 2nd Workshop on Hyperspectral Image and Signal Processing: Evolution in Remote Sensing*, pages 1–4. IEEE, 2010.
- [2] Andrew Boik and other. Improved audio mosaicing techniques for sound synthesis. Undergraduate senior thesis at Princeton University, 2013.
- [3] Ioan Buciu and Ioannis Pitas. Application of non-negative and local non negative matrix factorization to facial expression recognition. In *Proceedings of the 17th International Conference on Pattern Recognition, 2004. ICPR 2004.*, volume 1, pages 288–291. IEEE, 2004.
- [4] Gouyon F., Klapuri A., Dixon S., Alonso M., Tzanetakis G., Uhle C., and Cano P. An experimental comparison of audio tempo induction algorithms. *IEEE Transactions on Speech and Audio Processing*, 14(5), 2006.
- [5] Cédric Févotte, Nancy Bertin, and Jean-Louis Durrieu. Nonnegative matrix factorization with the itakura-saito divergence: With application to music analysis. *Neural computation*, 21(3):793–830, 2009.

- 136 [6] Derry Fitzgerald, Matt Cranitch, and Eugene Coyle. Shifted non-negative matrix factorisation
137 for sound source separation. In *IEEE/SP 13th Workshop on Statistical Signal Processing, 2005*,
138 pages 1132–1137. IEEE, 2005.
- 139 [7] Ryuji Hamamoto, Ken Takasawa, Hidenori Machino, Kazuma Kobayashi, Satoshi Takahashi,
140 Amina Bolatkan, Norio Shinkai, Akira Sakai, Rina Aoyama, Masayoshi Yamada, et al. Applica-
141 tion of non-negative matrix factorization in oncology: one approach for establishing precision
142 medicine. *Briefings in Bioinformatics*, 23(4):bbac246, 2022.
- 143 [8] Marko Helen and Tuomas Virtanen. Separation of drums from polyphonic music using non-
144 negative matrix factorization and support vector machine. In *2005 13th European Signal
145 Processing Conference*, pages 1–4. IEEE, 2005.
- 146 [9] Andre Holzapfel and Yannis Stylianou. Musical genre classification using nonnegative matrix
147 factorization-based features. *IEEE Transactions on Audio, Speech, and Language Processing*,
148 16(2):424–434, 2008.
- 149 [10] Daniel Lee and H Sebastian Seung. Algorithms for non-negative matrix factorization. *Advances
150 in neural information processing systems*, 13, 2000.
- 151 [11] Patricio López-Serrano, Christian Dittmar, Yigitcan Özer, and Meinard Müller. Nmf toolbox:
152 Music processing applications of nonnegative matrix factorization. In *Proceedings of the 22nd
153 International Conference on Digital Audio Effects (DAFx-19), Birmingham, UK*, pages 2–6,
154 2019.
- 155 [12] Ben Murrell, Thomas Weighill, Jan Buys, Robert Ketteringham, Sasha Moola, Gerdus Benade,
156 Lise Du Buisson, Daniel Kaliski, Tristan Hands, and Konrad Scheffler. Non-negative matrix fac-
157 torization for learning alignment-specific models of protein evolution. *PloS one*, 6(12):e28898,
158 2011.
- 159 [13] Ibrahim Selim, Arabi E Keshk, and Bassant M El Shourbugy. Galaxy image classification using
160 non-negative matrix factorization. *Int. J. Comput. Appl*, 137(5):4–8, 2016.
- 161 [14] [https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.
162 NMF.html](https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.NMF.html).